

EarthBridge: A Solution for 4th Multi-modal Aerial View Image Challenge Translation Track

Anonymous CVPR submission

Paper ID *****




Figure 1. Teaser: The proposed EarthBridge framework for multi-modal aerial-view image translation across Electro-Optical (EO), Infrared (IR), and Synthetic Aperture Radar (SAR) modalities. Our method accurately bridges the domain gap, preserving structural integrity while synthesizing high-fidelity target textures.

Abstract

001 Cross-modal image-to-image translation among Electro-
002 Optical (EO), Infrared (IR), and Synthetic Aperture Radar
003 (SAR) sensors is essential for comprehensive multi-modal
004 aerial-view analysis. However, translating between these
005 modalities is notoriously difficult due to their distinct elec-
006 tromagnetic signatures and geometric characteristics. This
007 paper presents **EarthBridge**, a high-fidelity translation
008 framework developed for the 4th Multi-modal Aerial View
009 Image Challenge – Translation (MAVIC-T). We explore two
010 distinct methodologies: **Diffusion Bridge Implicit Mod-**

els (DBIM), which we generalize using non-Markovian
bridge processes for high-quality deterministic sampling,
and **Contrastive Unpaired Translation (CUT)**, which uti-
lizes contrastive learning for structural consistency. Our
EarthBridge framework employs a channel-concatenated
UNet denoiser trained with Karras-weighted bridge scal-
ings and a specialized “booting noise” initialization to
handle the inherent ambiguity in cross-modal mappings.
We evaluate these methods across all four challenge tasks
(SAR $\rightarrow$ EO, SAR $\rightarrow$ RGB, SAR $\rightarrow$ IR, RGB $\rightarrow$ IR), achieving
superior spatial detail and spectral accuracy. Our solution
achieved a composite score of 0.38, securing the second po-

011
012
013
014
015
016
017
018
019
020
021
022

023 *sition on the MAVIC-T leaderboard.*

024 1. Introduction

025 The Multi-modal Aerial View Image Challenge for cross-
026 domain image translation [1, 19, 20] is a leading chal-
027 lenge held in conjunction with the Perception Beyond the
028 Visible Spectrum (PBVS) workshop from 2023 to 2025.
029 Data collected across multiple sensing modalities—such as
030 Electro-Optical (EO), Infrared (IR), and Synthetic Aper-
031 ture Radar (SAR)—provide complementary perspectives on
032 the same scene. However, data coverage limitations and
033 modality-specific constraints often hinder the full utiliza-
034 tion of these diverse sources. For instance, while SAR
035 sensors offer all-weather capabilities and can penetrate at-
036 mospheric obstructions, their imagery is complex to inter-
037 pret, and publicly available datasets are scarce. Conversely,
038 widely available EO imagery is limited by lighting condi-
039 tions and atmospheric visibility. Domain translation seeks
040 to bridge this gap by converting one modality into another,
041 enabling cross-modal data synthesis and enhancing sensor
042 interoperability. By focusing on translation tasks such
043 as RGB→IR, SAR→EO, SAR→IR, and SAR→RGB, the
044 MAVIC-T challenge aims to bolster existing data volume,
045 increase data diversity across regions, and serve as a plat-
046 form for advancing conditioned image synthesis techniques,
047 ultimately fostering the development of more adaptable and
048 robust models for multi-modal image analysis.

049 SAR images possess unique attributes that present chal-
050 lenges for both human observers and vision artificial intel-
051 ligence (AI) models to interpret, owing to their electromag-
052 netic (EM) characteristics [10–13, 21].

053 In this paper, we propose **EarthBridge**, a multi-method
054 framework for cross-modal aerial image translation. Our
055 primary approach is built on Diffusion Bridge Implicit
056 Models (DBIM) [43], which we generalize with non-
057 Markovian bridge processes to enable efficient determin-
058 istic sampling. In parallel, we explore Contrastive Un-
059 paired Translation (CUT) [24] as a powerful contrastive-
060 learning-based alternative. Unlike standard conditional dif-
061 fusion models that diffuse targets to pure noise and rely
062 on classifier-free guidance or simple concatenation for the
063 reverse path, EarthBridge implements a genuine stochastic
064 diffusion bridge. By directly constraining the forward
065 generative trajectory between the source modality x_T and
066 target modality x_0 , and parameterizing the UNet denoiser
067 to solve the bridge marginals under Karras-weighted scal-
068 ings [15], EarthBridge (DBIM) achieves high-fidelity trans-
069 lation across disparate modalities. We utilize pixel-level
070 modeling throughout both approaches to preserve the full
071 spectral and spatial resolution of the remote sensing data.

072 Our findings indicate that the deterministic DBIM sam-
073 pling significantly improves inference speed, achieving

high-quality results in as few as 5 sampling steps in
RGB→IR. Qualitatively, EarthBridge demonstrates a supe-
rior ability to reconstruct complex urban layouts and fine-
grained textures in challenging tasks such as SAR-to-RGB,
effectively bridging the domain gap. Quantitatively, our
solution achieves a composite task score of 0.269 for the
SAR→EO task and 0.38 overall, ultimately placing second
in the MAVIC-T challenge leaderboard.

Our main contributions are summarized as follows:

- We propose **EarthBridge**, a multi-method framework for paired cross-modal aerial image translation incorporating both Diffusion Bridge Implicit Models (DBIM) and Contrastive Unpaired Translation (CUT), providing a robust comparison between diffusion-based and contrastive-learning-based approaches.
- We develop a specialized training protocol using Karras-weighted bridge scalings and pixel-level modeling, optimized for the high dynamic range and resolution requirements of multimodal remote sensing data.
- We conduct comprehensive experiments on all four MAVIC-T tasks (SAR→EO, SAR→RGB, SAR→IR, RGB→IR), reporting competitive results and providing detailed quantitative and qualitative analysis.

2. Related Work

Translating across remote sensing (RS) modalities, such as SAR to optical (S2O), is critical for multisource data fusion. In this section, we provide a comprehensive literature review of image-to-image translation, diffusion bridge models, and cross-modality translation to complement the specialized discussion on remote sensing.

General Image-to-Image Translation. Image-to-image translation has been extensively studied in computer vision. Foundational works include Pix2Pix for paired conditional translation [14] and CycleGAN for unpaired translation [45]. Subsequent advances include contrastive learning for unpaired translation [24], zero-shot [25] and one-step translation with text-to-image models [26].

Diffusion Bridge Models. Unlike standard diffusion models that generate from noise, diffusion bridge models learn a stochastic process that directly connects source and target distributions, making them well-suited for paired image-to-image translation. Denoising Diffusion Bridge Models (DDBM) [44] define a variance-preserving bridge between data distributions. Related formulations include Brownian Bridge Diffusion Models (BBDM) [16] and its bidirectional extension BiBBDM [40], Image-to-Image Schrödinger Bridge (I^2SB) [17], Diffusion Schrödinger Bridge matching [34], Dual Diffusion Implicit Bridges [35], and Direct Diffusion Bridge with data consistency [2]. Re-

cent extensions include Consistency Diffusion Bridge [6], Diffusion Bridge Implicit Models [43], deterministic translation via denoising Brownian bridge with dual approximators [39], and adaptive domain shift for cross-modality translation [38].

Cross-Modality Translation. Image translation across diverse sensor modalities—grayscale, near-infrared (NIR), thermal colorization, and color transfer—has been addressed with probabilistic color correction [23], NIR imagery colorization [36], deep colorization [21, 41].

GAN-based RS Translation. Early attempts at translating between SAR and optical modalities leveraged Generative Adversarial Networks (GANs). For instance, Wang et al. [37] proposed a hierarchical latent feature approach to improve the fidelity of SAR-to-optical translation. Hu et al. [9] specifically addressed this task in fire-disturbed regions using GANs to reconstruct optical details from SAR signals, which is vital for post-fire damage assessment when clouds or smoke obscure optical sensors.

Diffusion Models for RS Translation. Reflecting the recent shift in generative modeling, diffusion models have been increasingly applied to RS translation tasks. Qin et al. [28] introduced a conditional diffusion model incorporating spatial-frequency refinement to capture both global structure and fine-grained textures in SAR-to-optical translation. To address the computational overhead of iterative sampling, Qin et al. [29] later developed an efficient end-to-end diffusion model capable of one-step SAR-to-optical translation, significantly reducing inference latency while maintaining high generation quality. Furthermore, Wang et al. [38] proposed adaptive domain shift mechanisms within diffusion frameworks to better handle the large distribution gaps inherent in cross-modality remote sensing datasets. These advancements underscore the growing role of stochastic bridge processes and diffusion-based refinement in enabling robust, high-resolution translation across heterogeneous sensor modalities.

3. Preliminary

Problem Formulation Let $\mathcal{X} \subset \mathbb{R}^{C_s \times H \times W}$ denote the source domain and $\mathcal{Y} \subset \mathbb{R}^{C_t \times H \times W}$ denote the target domain, where C_s and C_t are the number of channels in the source and target modalities, respectively. Given a set of spatially-aligned training pairs $\{(x_i, y_i)\}_{i=1}^N$ with $x_i \in \mathcal{X}$ and $y_i \in \mathcal{Y}$, the goal of cross-modal image-to-image translation is to learn a mapping $G : \mathcal{X} \rightarrow \mathcal{Y}$ that minimizes a composite distance between the generated image $\hat{y} = G(x)$ and the ground-truth target y . In the context

of the MAVIC-T challenge, four translation tasks are defined (see Tab. 1): SAR→EO (1 → 1 channel, 256×256), SAR→RGB (1 → 3, 1024×1024), SAR→IR (1 → 1, 1024×1024), and RGB→IR (3 → 1, 1024×1024).

Evaluation Metric The MAVIC-T challenge evaluates submissions using a composite score derived from three complementary metrics: LPIPS [42] (Learned Perceptual Image Patch Similarity, based on VGG-16) for perceptual fidelity, FID [8] (Fréchet Inception Distance, based on InceptionV3) for distributional similarity, and L1 norm for pixel-level structural accuracy. These are combined into a task score:

$$S_{\text{task}} = \frac{\frac{2}{\pi} \arctan(\text{FID}) + \text{LPIPS} + \text{L1}}{3}, \quad (1)$$

where the arctan normalization maps FID into the $[0, 1]$ range to balance its contribution against the other two metrics. The overall score is the average of the four task scores, with a penalty of +1 added for each unattempted domain, encouraging participants to develop generalizable solutions across all modality pairs.

4. EarthBridge: Multi-method Framework for RS Translation

We propose EarthBridge, a flexible framework that explores two distinct methodologies for cross-modal image translation: Diffusion Bridge Implicit Models (DBIM) and Contrastive Unpaired Translation (CUT).

4.1. Method 1: Diffusion Bridge Implicit Models (DBIM)

Diffusion models have recently revolutionized generative modeling, yet they are fundamentally designed to transport complex distributions to a standard Gaussian prior. Adapting them to translate between two arbitrary distributions, such as heterogenous remote sensing modalities, requires conditioning the diffusion process on a fixed endpoint. This is achieved via Doob’s h -transform [5, 30].

Stochastic Bridges via h -transform. Given a base diffusion process $dx_t = f(x_t, t)dt + g(t)dw_t$, the conditioned SDE that almost surely reaches a prescribed target $x_0 = y$ at $t = 0$ starting from a source $x_T = x$ at $t = T$ is given by:

$$dx_t = f(x_t, t)dt + g(t)dw_t + g(t)^2 \nabla_{x_t} \log p(x_0 = y | x_t)dt, \quad (2)$$

where $p(x_0 = y | x_t)$ is the transition density of the original process. When the endpoints are fixed, the resulting conditioned process is called a *diffusion bridge* [3, 7, 18,

27, 32, 33], which forms the foundational mechanism for transporting one state to another. In the context of cross-modal translation, x represents the source modality (e.g., SAR) and y the target modality (e.g., optical). This formulation ensures that the generative trajectory is directly constrained by both endpoints, providing a more principled framework for paired translation than standard conditional diffusion.

Denoising Diffusion Bridge Models. Denoising Diffusion Bridge Models (DDBM) [44] provide a tractable realization of Eq. 2. Under a variance-preserving (VP) schedule, the forward bridge process constructs intermediate states z_t as:

$$z_t = a_t x + b_t y + \sigma_t \epsilon, \quad \epsilon \sim \mathcal{N}(0, I), \quad (3)$$

where a_t, b_t, σ_t are schedule-dependent coefficients. While DDBM excels in high-fidelity generation, its inference typically relies on simulating Eq. 2 via computationally expensive numerical solvers (e.g., Heun or stochastic samplers), often requiring hundreds of function evaluations (NFE) for optimal results. In this work, we adopt the more efficient *implicit* formulation to accelerate this process.

4.2. EarthBridge: Non-Markovian DBIM

We present our DBIM approach which generalizes the diffusion bridge framework with non-Markovian bridge processes defined on discretized timesteps. Our solution leverages the deterministic sampling property of DBIM to achieve high-fidelity cross-modal translation with significantly fewer sampling steps compared to the vanilla DDBM. Fig. 2 shows an overview of our EarthBridge pipeline.

Non-Markovian Diffusion Bridges. In contrast to the continuous-time Markovian bridge processes in DDBM, we consider a family of non-Markovian diffusion bridges defined over a sequence of discretized timesteps $0 = t_0 < t_1 < \dots < t_N = T$, controlled by a variance parameter $\rho \in \mathbb{R}^{N-1}$. The UNet denoiser is trained to predict the bridge marginals governed by the schedule-dependent coefficients (a_t, b_t, c_t) , where the source image x is supplied to the network via channel concatenation. This concatenation is strictly an architectural mechanism to inject the mandatory endpoint parameter x into the SDE formulation, not an ad-hoc conditioning trick akin to standard image-to-image models. The joint probability distribution $q^{(\rho)}(z_{t_0:N-1} | x)$ is factorized through a chain of conditional distributions:

$$q^{(\rho)}(z_{t_0:N-1} | x) = q_0(z_{t_0}) \prod_{n=1}^{N-1} q^{(\rho)}(z_{t_n} | z_0, z_{t_{n+1}}, x), \quad (4)$$

where q_0 is the target modality distribution and each step $q^{(\rho)}$ is defined as:

$$\mathcal{N}(a_{t_n} x + b_{t_n} z_0 + \sqrt{c_{t_n}^2 - \rho_{t_n}^2} \eta_{n+1}, \rho_{t_n}^2 I), \quad (5)$$

with $\eta_{n+1} = (z_{t_{n+1}} - a_{t_{n+1}} x - b_{t_{n+1}} z_0) / c_{t_{n+1}}$. Here, a_t, b_t, c_t are known coefficients determined by the VP log-SNR schedule (Sec. 4.1). Under this construction, the marginal distributions $q^{(\rho)}(z_{t_n} | x)$ are preserved to be consistent with the continuous-time forward bridge (Eq. (3)). The variance parameter ρ controls the level of stochasticity in the sampling process, ranging from a purely Markovian stochastic process to a fully deterministic implicit model.

Reverse Generative Process and Sampling. The generative process is initialized at the source image $z_T = x$ and iteratively samples z_{t_n} from $z_{t_{n+1}}$ via a parameterized data predictor $D_\theta(z_{t_{n+1}}, t_{n+1}, x)$. The core updating rule is given by:

$$z_{t_n} = a_{t_n} x + b_{t_n} \hat{z}_0 + \sqrt{c_{t_n}^2 - \rho_{t_n}^2} \hat{\epsilon} + \rho_{t_n} \epsilon, \quad (6)$$


where $\hat{z}_0 = D_\theta(z_{t_{n+1}}, t_{n+1}, x)$ is the clean data prediction and $\hat{\epsilon}$ is the predicted noise at time t_{n+1} . Setting $\rho_n = 0$ for all steps yields the deterministic *Diffusion Bridge Implicit Model (DBIM)*, which can be viewed as an Euler discretization of a probability flow ODE specifically tailored for diffusion bridges.

Booting Noise. A critical refinement in our sampling procedure is the introduction of *booting noise* at the initial step ($t_N = T$). Since $c_{t_N} = 0$, the denominator in $\hat{\epsilon}$ leads to singularity. To address this, we force the initial step to be stochastic by choosing $\rho_{N-1} = c_{t_{N-1}}$, which injects a standard Gaussian booting noise ϵ_{boot} . This booting noise serves as the latent variable, encoding the one-to-many mapping ambiguity in RS translation (e.g., several plausible infrared images for a single SAR input).

Channel Handling for Mismatched Modalities. For tasks where source and target modalities have different channel counts (e.g., RGB→IR or SAR→RGB), we operate the bridge in a shared `model_channels` space (typically 3 for RGB-related tasks). Single-band modalities are expanded to `model_channels` by channel repetition before entering the UNet, ensuring the concatenation $[z_t, x]$ remains well-defined across all tasks.

4.3. Method 2: Contrastive Unpaired Translation (CUT)

In addition to our diffusion-based approach, we leverage Contrastive Unpaired Translation (CUT) [24] as an alternative method for cross-modal translation. Unlike diffusion



(a) UNet denoiser for 1024px tasks.






Figure 1: A schematic for Denoising Diffusion Bridge Models. DDBM uses a diffusion process guided by a drift adjustment (in blue) towards an endpoint $\mathbf{x}_T = \mathbf{y}$. They learn to reverse such a bridge process by matching the denoising bridge score (in orange), which allows one to reverse from $\mathbf{x}_T$ to $\mathbf{x}_0$ for any $\mathbf{x}_T = \mathbf{y} \sim q_{\text{data}}(\mathbf{y})$. The forward SDE process shown on the top is unidirectional while the probability flow ODE shown at the bottom is deterministic and bidirectional. White nodes are stochastic while grey nodes are deterministic.

(c) [TODO: Subfigure placeholder.]



(b) [TODO: Subfigure placeholder.]



(d) [TODO: Subfigure placeholder.]

Figure 2. UNet architecture of our DBIM for 1024px tasks (SAR→RGB, RGB→IR). We train at 512×512 and infer at full resolution. Source $\mathbf{x}$ is channel-concatenated with noisy sample $\mathbf{z}_t$ and processed by this denoiser. Bridge scalings a_t, b_t, c_t construct $\mathbf{z}_t$ during training; deterministic DBIM sampling generates $\hat{\mathbf{z}}_0$ during inference.

bridges, CUT utilizes contrastive learning to maintain structural correspondence. Our implementation uses a Patch-based Contrastive Learning (PCL) strategy to maximize the mutual information between corresponding patches of the source and generated target images.

5. Experiments

5.1. Dataset and Preprocessing

We use the MAVIC-T 2025 dataset from the challenge organizers, with imagery from UNICORN, USGS EROS HRO, and UMBRA across four regions: Bingham, Centerfield, Manhattan, and UC Davis. Tab. 1 summarizes the four translation tasks. SAR→EO operates at 256×256 ; the other three tasks use 1024×1024 . For SAR→IR and SAR→RGB, we spatially align SAR scenes with optical mosaics before resampling. We also create crop-augmented training sets for the 1024×1024 tasks to increase data diversity.

5.2. Architecture and Training Objectives

UNet Denoising Backbone Our denoiser D_θ is a UNet [31] that predicts the clean target from the noisy sam-

Table 1. Task specifications and training data for the MAVIC-T challenge. Crop augmentation uses 1024×1024 sliding windows with stride 512 (50% overlap); SAR→optical pairs are georeference-aligned.

Task	Resolution	Channels	Samples (Orig.)	Samples (Aug.)
SAR→EO	256×256	1→1	89,411	—
SAR→RGB	1024×1024	1→3	10,576	75,703
SAR→IR	1024×1024	1→1	10,576	74,317
RGB→IR	1024×1024	3→1	2,272	7,388

ple $\mathbf{z}_t$ conditioned on timestep t and source $\mathbf{x}$. We use task-specific scaling (128 base channels with resolution-dependent multipliers) and train with the Prodigy [22] optimizer. Detailed architectural parameters and source conditioning strategies are provided in Sec. 6.1.

Training Objective The primary training objective is a weighted mean squared error between the denoised prediction $D_\theta(\mathbf{z}_t, t, \mathbf{x})$ and the clean target $\mathbf{y}$:

$$\mathcal{L}_{\text{denoise}} = \mathbb{E}_{(\mathbf{x}, \mathbf{y}), t, \epsilon} [w(t) \cdot \|D_\theta(\mathbf{z}_t, t, \mathbf{x}) - \mathbf{y}\|_2^2], \quad (7)$$

Table 2. Main results on the MAVIC-T evaluation set. ↓ lower is better.

Track	Method	FID↓	LPIPS↓	L1↓	Score↓	NFE	Time
SAR→EO	DBIM 500	0.222	0.504	0.080	0.269	1000	0.42 s
SAR→RGB	DBIM 1000	0.878	0.636	0.214	0.576	2000	160 s
SAR→IR	CUT huge	0.645	0.603	0.145	0.464	1	0.47 s
RGB→IR	DBIM 5	0.357	0.151	0.090	0.200	10	1.5 s

where $w(t)$ is the Karras bridge weighting. The diffusion time t is sampled from a Karras inverse- ρ distribution. During inference, we utilize deterministic DBIM sampling for efficiency. Full technical details on the loss weighting, time-sampling distribution, and per-task training configurations are deferred to the appendix (Sec. 6).

Baseline: Contrastive Unpaired Translation (CUT) To assess the performance of our EarthBridge framework, we incorporate Contrastive Unpaired Translation (CUT) [24] as a primary baseline. Our implementation of CUT follows the original patch-based contrastive learning strategy, where patches from the generated target image are encouraged to be closer to the corresponding patches from the source image in a shared embedding space. We adapt the CUT architecture with ResNet-based generators and PatchGAN discriminators, scaled to match the resolution requirements of each task (e.g., medium tier for 256×256 and large/huge tiers for 1024×1024 as detailed in Sec. 6).

5.3. Quantitative Results

Tab. 2 presents the quantitative performance comparison between EarthBridge and the CUT baseline across the challenge tasks. The composite task score provides a balanced evaluation of perceptual, distributional, and structural fidelity. EarthBridge demonstrates consistent performance across both low- and high-resolution tasks.

5.4. Qualitative Results

Fig. 3 shows qualitative examples across all four tasks. EarthBridge preserves fine-grained structural information from the source modality while accurately synthesizing target textures, demonstrating consistent performance across both low- and high-resolution settings.

5.5. Failure Case and Challenging Scenario Analysis

[TODO: Add failure case analysis and challenging scenarios.]

6. Implementation Details

This section provides the comprehensive technical details of the EarthBridge architecture, training hyperparameters, and multi-modal refinement modules.

Table 3. DBIM (SAR→EO, SAR→RGB, RGB→IR).

Parameter	SAR→EO	SAR→RGB	RGB→IR
Base channels	64	96	128
Channel multipliers	(1,2,3,4)	(1,1,2,2,4,4)	
Attention resolutions	N/A	32, 16, 8	
Training resolution	256^2	512^2	
Epochs	55	24	86

Table 4. CUT (SAR→IR). ResNet-9block, 4-layer PatchGAN.

Parameter	SAR→IR
ngf / ndf	256
Generator	ResNet-9block
Discriminator	4-layer PatchGAN
Train res.	512^2
Epochs	11

6.1. Architecture, Sampling, and Scheduler Settings

The EarthBridge architecture and sampling configurations are designed for high-fidelity cross-modal translation with efficiency. We utilize a UNet2DModel backbone (ADM-style [4]) with per-task architectural parameters: number of channels, number of residual blocks, attention resolutions, channel multipliers, resolution, and training settings (epochs, augmentations) as detailed in Tab. 3. Timestep information is injected via sinusoidal embeddings and Adaptive Group Normalization (AdaGN). Source image conditioning is implemented via channel-wise concatenation (concat mode) $[z_t, x]$.

Sampling is performed using the deterministic DBIM formulation [43] with task-dependent inference steps (see Tab. 2) and a Gaussian booting noise ϵ_{boot} at $t = T$. The bridge process follows a variance-preserving (VP) log-SNR schedule with $\beta_d = 2.0$ and $\beta_{\min} = 0.1$. For training, we sample diffusion time from a Karras inverse- ρ distribution with $\rho = 7$. Per-task architectural and training configurations are summarized in Tab. 3.

The CUT baseline follows the standard configuration with ResNet-based generators and PatchGAN discriminators. We use the *medium* tier for the 256×256 SAR→EO task and the *huge* tier for the 1024×1024 tasks to ensure sufficient capacity for high-resolution synthesis. The training hyperparameters for CUT, including the number of filters and residual blocks, are detailed in Tab. 4.

We train EarthBridge using the Prodigy adaptive optimizer across all translation tasks. The training objective




Figure 3. Qualitative results across all four MAVIC-T translation tasks. Each row corresponds to one task (SAR→EO, SAR→RGB, SAR→IR, RGB→IR). Columns show: source image, EarthBridge output, and ground-truth target. Our model preserves structural layout from the source while synthesizing faithful target-modality textures across both 256×256 and 1024×1024 resolutions.

401 follows the weighted Karras bridge loss with task-specific
402 batch sizes and augmentations.

403 7. Lessons and Insights

404 We summarize key lessons and insights from developing
405 EarthBridge for the MAVIC-T translation track.

Insight 1: Diffusion Bridge vs. Standard Diffusion

The diffusion bridge directly builds a mapping from the source domain to the target domain, whereas standard diffusion starts from a noise prior (typically Gaussian noise) and requires a conditional generation strategy to accomplish this task. We hypothesize that the diffusion bridge is more efficient for training and achieves better results for cross-modal translation. However, standard diffusion can be more flexible for further conditional development, e.g., leveraging pre-trained models for a good prior. There is a trade-off between these approaches; we encourage readers to further explore both directions.

406

Insight 2: Resolution Scaling and Multi-Stage Curriculum

For 1024px modeling, we train on 512px only and directly infer on 1024px, which yields strong results while saving training compute budget. If time permits, we may further perform a quick fine-tuning on 1024px directly—we term this approach *multi-stage resolution curriculum learning*.

ther optimization for real-time remote sensing applications.
[TODO: Add final ranking and summary of key findings.]

428
429
430

Insight 3: Task Difficulty and Input Modality

We found that SAR→EO and RGB→IR are substantially easier than the other two tasks. For RGB→IR, the three-channel input provides sufficient information for the translation. For SAR→EO, the relative ease may stem from both modalities originating from the same sensor, reducing the domain gap.

Insight 4: Inference Steps and Metric Trade-offs

Using the DBIM ODE solver, few-step inference yields visually appealing results. For RGB→IR, even a single-step generation is sufficient, and increasing steps tends to hallucinate details. In practice, however, many more inference steps greatly improve FID and LPIPS in the competition metrics, even when L1 remains unchanged—highlighting a trade-off between perceptual quality, metric optimization, and inference efficiency.

8. Lessons and Insights

We summarize key lessons and insights from developing EarthBridge for the MAVIC-T translation track.

9. Conclusion

In this paper, we presented EarthBridge, a novel solution for the 4th Multi-modal Aerial View Image Challenge – Translation (MAVIC-T). By leveraging Diffusion Bridge Implicit Models (DBIM), we generalized the diffusion bridge framework with non-Markovian processes, enabling efficient and deterministic cross-modal translation. Our approach effectively handles disparate sensor characteristics across EO, IR, and SAR modalities through pixel-level modeling and a specialized training protocol featuring Karras-weighted bridge scalings. Quantitative and qualitative evaluations across four translation tasks demonstrate EarthBridge’s superior ability to preserve spatial structure while generating high-fidelity target textures. Future work will explore the integration of adaptive domain shift mechanisms and fur-

References

- [1] Dylan Bowald, Justice Wheelwright, Oliver Nina, Angel Sappa, Riad Hammoud, Erik Blasch, and Nathan Inkawhich. 3rd Multi-modal Aerial View Image Challenge: Sensor Domain Translation - PBVS 2025. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 4679–4688, 2025. 2
- [2] Hyungjin Chung, Jeongsol Kim, and Jong Chul Ye. Direct Diffusion Bridge using Data Consistency for Inverse Problems. In *Neural Information Processing Systems*, pages 7158–7169. Curran Associates, Inc., 2023. 2
- [3] Bernard Delyon and Ying Hu. Simulation of conditioned diffusion and application to parameter estimation. *Stochastic Processes and Their Applications*, 116(11):1660–1675, 2006. 3
- [4] Prafulla Dhariwal and Alexander Nichol. Diffusion Models Beat GANs on Image Synthesis. In *Advances in Neural Information Processing Systems*, pages 8780–8794. Curran Associates, Inc., 2021. 6
- [5] Joseph L. Doob and J. I. Doob. *Classical Potential Theory and Its Probabilistic Counterpart*. Springer, 1984. 3
- [6] Guande He, Kaiwen Zheng, Jianfei Chen, Fan Bao, and Jun Zhu. Consistency Diffusion Bridge Models. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024. 3
- [7] Jeremy Heng, Valentin De Bortoli, Arnaud Doucet, and James Thornton. Simulating diffusion bridges with score matching. *arXiv preprint arXiv:2111.07243*, 2021. 3
- [8] Martin Heusel, Hubert Ramsauer, Thomas Unterthiner, Bernhard Nessler, and Sepp Hochreiter. GANs Trained by a Two Time-Scale Update Rule Converge to a Local Nash Equilibrium. In *Advances in Neural Information Processing Systems*. Curran Associates, Inc., 2017. 3
- [9] Xikun Hu, Puzhao Zhang, and Yifang Ban. Gan-based SAR to Optical Image Translation in Fire-Disturbed Regions. In *IGARSS 2022 - 2022 IEEE International Geoscience and Remote Sensing Symposium*, pages 1456–1459, 2022. 3
- [10] Zhongling Huang, Mihai Datcu, Zongxu Pan, and Bin Lei. Deep SAR-Net: Learning objects from signals. *ISPRS Journal of Photogrammetry and Remote Sensing*, 161:179–193, 2020. 2
- [11] Zhongling Huang, Mihai Datcu, Zongxu Pan, Xiaolan Qiu, and Bin Lei. HDEC-TFA: An Unsupervised Learning Approach for Discovering Physical Scattering Properties of Single-Polarized SAR Image. *IEEE Transactions on Geoscience and Remote Sensing*, 59(4):3054–3071, 2021.
- [12] Zhongling Huang, Corneliu Octavian Dumitru, Zongxu Pan, Bin Lei, and Mihai Datcu. Classification of Large-Scale High-Resolution SAR Images With Deep Transfer Learning. *IEEE Geoscience and Remote Sensing Letters*, 18(1):107–111, 2021. 484, 485, 486, 487, 488
- [13] Zhongling Huang, Xiwen Yao, Ying Liu, Corneliu Octavian Dumitru, Mihai Datcu, and Junwei Han. Physically explainable CNN for SAR image classification. *ISPRS Journal of Photogrammetry and Remote Sensing*, 190:25–37, 2022. 2 489, 490, 491, 492, 493
- [14] Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, and Alexei A. Efros. Image-To-Image Translation With Conditional Adversarial Networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 1125–1134, 2017. 2 494, 495, 496, 497, 498
- [15] Tero Karras, Miika Aittala, Timo Aila, and Samuli Laine. Elucidating the design space of diffusion-based generative models. In *Adv. Neural Inform. Process. Syst.*, 2022. 2 499, 500, 501, 502
- [16] Bo Li, Kaitao Xue, Bin Liu, and Yu-Kun Lai. BBDM: Image-to-image translation with brownian bridge diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 1952–1961, 2023. 2 503, 504, 505, 506, 507
- [17] Guan-Hong Liu, Arash Vahdat, De-An Huang, Evangelos Theodorou, Weili Nie, and Anima Anandkumar. I²sb: Image-to-Image Schrödinger Bridge. In *The Eleventh International Conference on Learning Representations*, 2023. 2 508, 509, 510, 511, 512
- [18] Xingchao Liu, Lemeng Wu, Mao Ye, and Qiang Liu. Let us build bridges: Understanding and extending diffusion generative models. *arXiv preprint arXiv:2208.14699*, 2022. 3 513, 514, 515, 516
- [19] Spencer Low, Oliver Nina, Angel D. Sappa, Erik Blasch, and Nathan Inkawhich. Multi-modal Aerial View Image Challenge: Translation from Synthetic Aperture Radar to Electro-Optical Domain Results - PBVS 2023. In *2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 515–523, 2023. 2 517, 518, 519, 520, 521, 522, 523
- [20] Spencer Low, Oliver Nina, Dylan Bowald, Angel D. Sappa, Nathan Inkawhich, and Peter Bruns. Multi-modal Aerial View Image Challenge: Sensor Domain Translation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 3096–3104, 2024. 2 524, 525, 526, 527, 528, 529
- [21] Uttam K Majumder, Erik P Blasch, and David A Garren. *Deep learning for radar and communications automatic target recognition*. Artech House, 2020. 2, 3 530, 531, 532, 533
- [22] Konstantin Mishchenko and Aaron Defazio. Prodigy: An Expediently Adaptive Parameter-Free Learner. In *Proceedings of the 41st International Conference* 534, 535, 536

- on *Machine Learning*, pages 35779–35804. PMLR, 2024. 5
- [23] Miguel Oliveira, Angel Domingo Sappa, and Vitor Santos. A Probabilistic Approach for Color Correction in Image Mosaicking Applications. *IEEE Transactions on Image Processing*, 24(2):508–523, 2015. 3
- [24] Taesung Park, Alexei A. Efros, Richard Zhang, and Jun-Yan Zhu. Contrastive Learning for Unpaired Image-to-Image Translation. In *Computer Vision—ECCV 2020: 16th European Conference, Glasgow, UK, August 23–28, 2020, Proceedings, Part IX 16*, pages 319–345, Cham, 2020. Springer International Publishing. 2, 4, 6
- [25] Gaurav Parmar, Krishna Kumar Singh, Richard Zhang, Yijun Li, Jingwan Lu, and Jun-Yan Zhu. Zero-shot Image-to-Image Translation. In *ACM SIGGRAPH 2023 Conference Proceedings*, pages 1–11, New York, NY, USA, 2023. Association for Computing Machinery. 2
- [26] Gaurav Parmar, Taesung Park, Srinivasa Narasimhan, and Jun-Yan Zhu. One-Step Image Translation with Text-to-Image Models, 2024. 2
- [27] Stefano Peluchetti. Non-denoising forward-time diffusions. *arXiv preprint arXiv:2109.11704*, 2021. 4
- [28] Jiang Qin, Kai Wang, Bin Zou, Lamei Zhang, and Joost van de Weijer. Conditional Diffusion Model With Spatial-Frequency Refinement for SAR-to-Optical Image Translation. *IEEE Transactions on Geoscience and Remote Sensing*, 62:1–14, 2024. 3
- [29] Jiang Qin, Bin Zou, Haolin Li, and Lamei Zhang. Efficient End-to-End Diffusion Model for One-Step SAR-to-Optical Translation. *IEEE Geoscience and Remote Sensing Letters*, 22:1–5, 2025. 3
- [30] L. Chris G. Rogers and David Williams. *Diffusions, Markov Processes and Martingales: Volume 2, Itô Calculus*. Cambridge University Press, 2000. 3
- [31] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-Net: Convolutional Networks for Biomedical Image Segmentation. In *Medical Image Computing and Computer-Assisted Intervention – MICCAI 2015*, pages 234–241, Cham, 2015. Springer. 5
- [32] Simo Särkkä and Arno Solin. *Applied Stochastic Differential Equations*. Cambridge University Press, 2019. 4
- [33] Moritz Schauer, Frank Van Der Meulen, and Harry Van Zanten. Guided proposals for simulating multi-dimensional diffusion bridges. *Bayesian Analysis*, 12(2):443–471, 2017. 4
- [34] Yuyang Shi, Valentin De Bortoli, Andrew Campbell, and Arnaud Doucet. Diffusion Schrödinger Bridge Matching. In *Neural Information Processing Systems*, pages 62183–62223. Curran Associates, Inc., 2023. 2
- [35] Xuan Su, Jiaming Song, Chenlin Meng, and Stefano Ermon. Dual Diffusion Implicit Bridges for Image-to-Image Translation. In *The Eleventh International Conference on Learning Representations*, 2023. 2
- [36] Patricia L. Suárez, Angel D. Sappa, Boris X. Vintimilla, and Riad I. Hammoud. Near InfraRed Imagery Colorization. In *2018 25th IEEE International Conference on Image Processing (ICIP)*, pages 2237–2241, 2018. 3
- [37] Haixia Wang, Zhigang Zhang, Zhanyi Hu, and Qiulei Dong. SAR-to-Optical Image Translation With Hierarchical Latent Features. *IEEE Transactions on Geoscience and Remote Sensing*, 60:1–12, 2022. 3
- [38] Zihao Wang, Yuzhou Chen, and Shaogang Ren. Adaptive Domain Shift in Diffusion Models for Cross-Modality Image Translation. In *The Fourteenth International Conference on Learning Representations*, 2026. 3
- [39] Bohan Xiao, Peiyong Wang, Qisheng He, and Ming Dong. Deterministic Image-to-Image Translation via Denoising Brownian Bridge Models with Dual Approximators. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 28232–28241, 2025. 3
- [40] Kaitao Xue, Bo Li, Ziyi Liu, Zhifen He, Bin Liu, Congxuan Zhang, and Yu-Kun Lai. BiBDDM: Bidirectional Image Translation With Brownian Bridge Diffusion Models. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 47(11):10546–10559, 2025. 2
- [41] Richard Zhang, Phillip Isola, and Alexei A. Efros. Colorful Image Colorization. In *Computer Vision – ECCV 2016*, pages 649–666, Cham, 2016. Springer International Publishing. 3
- [42] Richard Zhang, Phillip Isola, Alexei A. Efros, Eli Shechtman, and Oliver Wang. The Unreasonable Effectiveness of Deep Features as a Perceptual Metric. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 586–595, 2018. 3
- [43] Kaiwen Zheng, Guande He, Jianfei Chen, Fan Bao, and Jun Zhu. Diffusion Bridge Implicit Models. In *The Thirteenth International Conference on Learning Representations*, 2025. 2, 3, 6
- [44] Linqi Zhou, Aaron Lou, Samar Khanna, and Stefano Ermon. Denoising Diffusion Bridge Models. In *The Twelfth International Conference on Learning Representations*, 2024. 2, 4
- [45] Jun-Yan Zhu, Taesung Park, Phillip Isola, and Alexei A. Efros. Unpaired Image-To-Image Translation Using Cycle-Consistent Adversarial Networks. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 2223–2232, 2017. 2

EarthBridge: A Solution for 4th Multi-modal Aerial View Image Challenge Translation Track

Supplementary Material