

EarthBridge: A Solution for the 4th Multi-modal Aerial View Image Challenge – Translation Track

Team: EarthBridge **CodaBench:** bilisakura **Code:** github.com/Bili-Sakura/EarthBridge-Preview

Zhenyuan Chen[†]
Zhejiang University

bili_sakura@zju.edu.cn

Yuanshen Guan[†]
University of Science and Technology of China
guanyans@mail.ustc.edu.cn

Feng Zhang
Zhejiang University
zfcarnation@zju.edu.cn

Abstract

We present **EarthBridge**, our solution for the 4th MAVIC-T challenge. We apply Diffusion Bridge Implicit Models (DBIM) [6] to three translation tasks (SAR→EO, SAR→RGB, RGB→IR) and Contrastive Unpaired Translation (CUT) for SAR→IR. Our solution achieves the second position overall. We report composite scores of 0.269, 0.576, 0.464, and 0.200 on SAR→EO, SAR→RGB, SAR→IR, and RGB→IR respectively, with an average of 0.38.

1. Method

1.1. Method 1: Diffusion Bridge Implicit Models (DBIM)

1.1.1. Bridge Design and Training

Stochastic Bridges. Given a base diffusion process $dx_t = f(x_t, t)dt + g(t)dw_t$, the bridge conditioned to reach target $x_0 = \mathbf{y}$ starting from source $x_T = \mathbf{x}$ is:

$$dx_t = f(x_t, t)dt + g(t)dw_t + g(t)^2 \nabla_{x_t} \log p(x_0 = \mathbf{y} | x_t) dt. \quad (1)$$

Under a variance-preserving (VP) schedule, intermediate states are:

$$z_t = a_t \mathbf{x} + b_t \mathbf{y} + \sigma_t \epsilon, \quad \epsilon \sim \mathcal{N}(0, I), \quad (2)$$

where $\mathbf{x}$ is the source image, $\mathbf{y}$ is the target image, and a_t, b_t, σ_t are schedule coefficients [7].

Non-Markovian Generalization and Booting Noise.

We extend the Markovian bridge to a family of non-Markovian processes parametrized by variance $\rho \in \mathbb{R}^{N-1}$.

The reverse update rule is:

$$z_{t_n} = a_{t_n} \mathbf{x} + b_{t_n} \hat{z}_0 + \sqrt{c_{t_n}^2 - \rho_n^2} \hat{\epsilon} + \rho_n \epsilon, \quad (3)$$

where $\hat{z}_0 = D_\theta(z_{t_{n+1}}, t_{n+1}, \mathbf{x})$ is the UNet denoiser prediction. Setting $\rho_n = 0$ gives the deterministic DBIM [6], analogous to an Euler discretization of a probability-flow ODE over the bridge. Since $c_{t_N} = 0$ leads to a singularity at initialization, we force the first step to be stochastic by setting $\rho_{N-1} = c_{t_{N-1}}$, injecting Gaussian booting noise that encodes the one-to-many mapping ambiguity.

Training Objective. The model is trained with Karras-weighted bridge loss:

$$\mathcal{L} = \mathbb{E}_{(\mathbf{x}, \mathbf{y}), t, \epsilon} [w(t) \cdot \|D_\theta(z_t, t, \mathbf{x}) - \mathbf{y}\|_2^2], \quad (4)$$

where $w(t)$ follows Karras bridge weighting [2] and t is sampled from a Karras inverse- ρ distribution.

1.1.2. Architecture and Channel Handling

Our denoiser D_θ is a UNet [5] that receives the channel-wise concatenation $[z_t, \mathbf{x}]$ as input. Timestep information is injected via sinusoidal embeddings and Adaptive Group Normalization (AdaGN, ADM-style [1]). The architecture overview is shown in Fig. 1.

For tasks where source and target modalities have different channel counts (e.g., RGB→IR or SAR→RGB), we operate the bridge in a shared channel space (typically 3 for RGB-related tasks). Single-band modalities are expanded by channel repetition before entering the UNet, ensuring the 1-to-1 mapping. **Method 2: Contrastive Unpaired Translation (CUT)**

For the SAR→IR task, we employ CUT [4] as our final method. CUT uses patch-based contrastive learning




Figure 1. UNet architecture of our DBIM for 1024px tasks (SAR→RGB, RGB→IR). We train at 512×512 and infer at full resolution. Source x is channel-concatenated with noisy sample z_t and processed by this denoiser. Bridge scalings a_t, b_t, c_t construct z_t during training; deterministic DBIM sampling generates $\hat{z}_0$ during inference.

to maintain structural correspondence between source and generated target, enabling effective translation even without pixel-aligned pairs. We use the *huge* tier (256 ResNet filters, 4-layer PatchGAN) to handle the 1024×1024 resolution. For channel handling: the single-band SAR input and single-band IR output are processed directly without modality expansion.

2. Experiments

2.1. Dataset

SAR→EO uses 256×256 imagery with 89,411 pairs. The other three tasks use 1024×1024 : SAR→RGB (10,576 orig., 75,703 aug.), SAR→IR (10,576 orig., 74,317 aug.), RGB→IR (2,272 orig., 7,388 aug.). For SAR→IR and SAR→RGB, we spatially align SAR scenes with optical mosaics before resampling. For 1024×1024 tasks, we augment training data with sliding-window crops (stride 512).

2.2. Training Details

All DBIM models are trained with the Prodigy adaptive optimizer [3]. For 1024×1024 tasks, we train at 512×512 and infer at full resolution (resolution curriculum). The bridge follows a VP log-SNR schedule with $\beta_d = 2.0$, $\beta_{\min} = 0.1$. We use DBIM for SAR→EO, SAR→RGB, and RGB→IR; CUT for SAR→IR only. CUT (huge tier) is trained with Prodigy, batch size 2 per GPU across 8 GPUs, for 11 epochs, with NCE temperature 0.1 and $\lambda_{\text{GAN}} = 0.5$. Configurations are given in Tabs. 1 and 2.

2.3. Hardware and Software

Hardware. Eight NVIDIA H200 GPUs; x86_64 multi-core CPU (AMD EPYC / Intel Xeon); 512 GB+ RAM.

Table 1. DBIM (SAR→EO, SAR→RGB, RGB→IR).

Parameter	SAR→EO	SAR→RGB	RGB→IR
Base ch.	64	96	128
Ch. mults	(1,2,3,4)	(1,1,2,2,4,4)	
Attn res.	N/A	32, 16, 8	
Train res.	256^2	512^2	
Epochs	55	24	86

Table 2. CUT (SAR→IR). ResNet-9block, 4-layer PatchGAN.

Parameter	SAR→IR
ngf / ndf	256
Generator	ResNet-9block
Discriminator	4-layer PatchGAN
Train res.	512^2
Epochs	11

Software. Python 3.10+; PyTorch 2.8; HuggingFace Diffusers 0.36; Transformers; Accelerate; Rasterio (geospatial). Code: github.com/Bili-Sakura/EarthBridge-Preview.

2.4. Quantitative Results

Tab. 3 reports the performance of our **EarthBridge** final submission. NFE denotes number of function evaluations (for DBIM, NFE = steps $\times$ 2; for CUT, NFE = 1).

Table 3. Results on the MAVIC-T evaluation set. ↓ lower is better.

Track	Method	FID↓	LPIPS↓	L1↓	Score↓	NFE	Time
SAR→EO	DBIM (500)	0.222	0.504	0.080	0.269	1000	0.42 s
SAR→RGB	DBIM (1000)	0.878	0.636	0.214	0.576	2000	160 s
SAR→IR	CUT (huge)	0.645	0.603	0.145	0.464	1	0.47 s
RGB→IR	DBIM (5)	0.357	0.151	0.090	0.200	10	1.5 s

Key Observations. (1) SAR→EO and RGB→IR are extremely easy to model; we achieve high performance on both test and validation sets, which aligns with peer rankings in previous years. (2) SAR→RGB and SAR→IR are quite difficult, also consistent with prior competition results. Given the difficulty of SAR→IR, we switched from DBIM to the GAN-based CUT. This yielded better results: diffusion bridges for this hard task tend to produce over-smooth, blurry outputs with low training loss that nonetheless deviate from the actual target appearance; CUT’s discriminator addresses this. (3) Although inference requires full 1024×1024 resolution, we find that training our UNet at 512×512 effectively extends to 1024×1024 inference without modification. (4) Our solution is computationally heavy: we employ a large model (about 200M parameters) and train on a node cluster with eight NVIDIA H200 GPUs for approximately ten hours.

References

- [1] Prafulla Dhariwal and Alexander Nichol. Diffusion Models Beat GANs on Image Synthesis. In *Advances in Neural Information Processing Systems*, pages 8780–8794. Curran Associates, Inc., 2021. [1](#)
- [2] Tero Karras, Miika Aittala, Timo Aila, and Samuli Laine. Elucidating the design space of diffusion-based generative models. In *Adv. Neural Inform. Process. Syst.*, 2022. [1](#)
- [3] Konstantin Mishchenko and Aaron Defazio. Prodigy: An Expeditiously Adaptive Parameter-Free Learner. In *Proceedings of the 41st International Conference on Machine Learning*, pages 35779–35804. PMLR, 2024. [2](#)
- [4] Taesung Park, Alexei A. Efros, Richard Zhang, and Jun-Yan Zhu. Contrastive Learning for Unpaired Image-to-Image Translation. In *Computer Vision—ECCV 2020: 16th European Conference, Glasgow, UK, August 23–28, 2020, Proceedings, Part IX 16*, pages 319–345, Cham, 2020. Springer International Publishing. [1](#)
- [5] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-Net: Convolutional Networks for Biomedical Image Segmentation. In *Medical Image Computing and Computer-Assisted Intervention – MICCAI 2015*, pages 234–241, Cham, 2015. Springer. [1](#)
- [6] Kaiwen Zheng, Guande He, Jianfei Chen, Fan Bao, and Jun Zhu. Diffusion Bridge Implicit Models. In *The Thirteenth International Conference on Learning Representations*, 2025. [1](#)
- [7] Linqi Zhou, Aaron Lou, Samar Khanna, and Stefano Ermon. Denoising Diffusion Bridge Models. In *The Twelfth International Conference on Learning Representations*, 2024. [1](#)